Dynamic Financial Portfolio Optimization Using Temporal Convolutional Networks and Real-Time Data Analysis

Abhijith E
Department of Computer Science
Christ University
Bangalore, India
abhijith.e@msam.christuniversity.in

Gobi Ramasamy
Department of Computer Science
Christ University
Bangalore, India
gobi.r@christuniversity.in

Saravanan KN
Department of Computer Science
Christ University
Bangalore, India
saravanan.kn@christuniversity.in

***Abstract*—This paper presents an integrated framework for AI-driven portfolio optimization combining temporal convolutional networks (TCNs) with conditional value-at-risk (CVaR) minimization. Our system processes real-time market data through an automated pipeline implementing volatility-adjusted feature engineering and walk-forward validation. The architecture employs dilated causal convolutions for temporal pattern extraction combined with Ledoit-Wolf shrinkage covariance estimation for robust portfolio optimization. Experimental results demonstrate an 18.7% annualized return with 22.3% volatility, outperforming traditional mean-variance optimization by 14.2% in risk-adjusted returns. The implementation addresses key challenges in numerical stability and computational efficiency through eigenvalue clamping and gradient checkpointing.**

***Index Terms*—Portfolio optimization, temporal convolutional networks, real-time data analysis, risk management, algorithmic trading**

I. INTRODUCTION

Modern financial markets are highly dynamic, exhibiting rapid fluctuations, unexpected volatility, and complex interdependencies among various assets. Traditional portfolio management techniques, while effective in stable conditions, often fail to adapt quickly to changing market environments. The limitations of classical optimization methods become evident in volatile markets where assumptions of stability and predictability no longer hold. This research aims to address these challenges by introducing an advanced framework that integrates machine learning techniques with robust statistical models to enhance the adaptability and efficiency of portfolio management.Traditional optimization methods [1] face three fundamental challenges:

A. Non-stationarity in asset return distributions

- Traditional financial models often assume that asset returns follow a stationary distribution, meaning that statistical properties such as mean and variance remain constant over time. However, real-world financial markets are inherently non-stationary due to factors such as macroeconomic shifts, geopolitical events, and evolving investor sentiment.
- Failure to account for non-stationarity can lead to misleading risk estimations and suboptimal portfolio allocations.
- To address this challenge, we employ Temporal Convolutional Networks (TCNs) with dilated causal convolutions, which enable the detection of multi-scale patterns in financial time-series data. By capturing both short-term and long-term dependencies while preserving the chronological order of events, this approach ensures a more accurate assessment of asset return behavior.

B. High-dimensional covariance matrix estimation

- Estimating the covariance matrix of asset returns is crucial for portfolio optimization, as it defines the relationships between different assets and their collective risk exposure. However, when dealing with large portfolios containing numerous assets, traditional estimation methods struggle due to the curse of dimensionality and the limited availability of historical data.
- Conventional techniques, such as sample covariance estimation, often result in unstable and noisy covariance matrices, leading to unreliable portfolio allocations.
- Our research proposes a hybrid covariance estimation technique that combines Ledoit-Wolf shrinkage with eigenvalue regularization. The Ledoit-Wolf shrinkage method improves estimation accuracy by reducing noise and enhancing the stability of covariance matrices, while eigenvalue regularization ensures that the estimated matrix remains positive semi-definite, improving its reliability in risk assessments.

C. Latency in traditional risk metric calculations

- Expected Shortfall (ES) and Value-at-Risk (VaR) are among the basic risk measures of risk management. Conventional calculation methods for these measures, however, incur latency, inhibiting real-time risk estimation and portfolio rebalancing.

- Market conditions constantly shift, and procrastination in measuring risk may result in lost opportunities or heightened risk exposure to undesirable market movement.
- To resolve this problem, we create an automated data pipeline based on volatility-adjusted feature engineering. The pipeline adapts input features in real-time to changing market conditions and updates risk measures in real-time. With real-time data and adaptive learning models, our system makes portfolio decisions based on the most recent and most pertinent data.

Our answer addresses these problems with three breakthroughs:

1) Temporal convolutional networks with dilated causal convolutions for multi-scale pattern recognition.
2) Hybrid covariance estimation combining Ledoit-Wolf shrinkage with eigenvalue regularization
3) Automated data pipeline implementing volatility-adjusted feature engineering

II. Related Work

The application of Artificial Intelligence (AI) and Machine Learning (ML) in financial markets has seen rapid advancements, particularly in stock price prediction, portfolio management, and risk assessment. Recent research integrates sentiment analysis, reinforcement learning, and deep learning techniques to enhance predictive accuracy and decision-making capabilities. Early research primarily relied on statistical methods like Support Vector Regression (SVR) and Decision Trees for stock price prediction, providing interpretable results but struggling with real-time adaptability and complex, high-dimensional data. With the rise of social media and financial news, sentiment analysis has become an essential component of stock prediction, with studies demonstrating the influence of Twitter sentiment and microblogging data on stock movements. Models like Dirichlet Mixture, Kalman Filters, CNNs, and LSTMs have improved prediction accuracy by integrating textual sentiment with numerical stock data. Deep learning-based forecasting highlights the superiority of CNNs and LSTMs in time-series forecasting, learning patterns from historical and real-time data. For example, StockSentiWordNet (SSWN) combined sentiment analysis with extreme learning machines (ELM) and RNNs to achieve an 86 Percentage accuracy in stock trend prediction. Traditional portfolio management approaches like ”Follow-the-Winner” and ”Meta-Learning” struggle with market adaptability. Reinforcement Learning (RL) frameworks, such as Proximal Policy Optimization (PPO) and Advantage Actor-Critic (A2C), have demonstrated enhanced decision-making in portfolio rebalancing. Research shows that CNN-based ensemble models outperform conventional techniques in asset allocation and risk control. AI-driven risk assessment techniques have revolutionized financial decision-making, with unsupervised learning methods like K-Means clustering helping categorize stocks based on market characteristics. Supervised learning models, including Decision Trees and Random Forests, predict default risk and market volatility, while deep learning-based hedging strategies show promising results in mitigating credit and market risks. HFT strategies require precise trade execution within milliseconds, and reinforcement learning algorithms optimize order placement and execution costs using bid-ask spreads and transaction volumes. Studies on RL-based trade execution highlight their ability to outperform traditional static strategies. The use of ML in market microstructure data analysis enables better price movement predictions and order routing, where features like order book imbalance and trade flow statistics enhance profitability in dark pool trading. Market risk models leveraging deep learning have shown higher accuracy in predicting financial crises, while AI-based credit scoring systems outperform traditional models by integrating alternative data sources like purchase histories and online behavior. Reinforcement learning is being explored for stress testing and risk mitigation. AI’s role in regulatory technology (RegTech) focuses on automating compliance checks using NLP and deep learning, with IBM Watson utilized for regulatory audits, demonstrating improvements in fraud detection and anti-money laundering (AML) processes. AI models, particularly deep learning, face challenges in interpretability, with ”black-box” decision-making raising concerns in finance, prompting the development of Explainable AI (XAI) frameworks. The ethical implications of AI-driven financial decision-making include bias, automation-induced job displacement, and market manipulation risks, necessitating responsible AI initiatives. Emerging trends like quantum computing and blockchain integration are expected to further revolutionize AI in finance, while hybrid AI models combining classical statistical techniques with deep learning offer new opportunities for improving financial decision-making. The integration of AI and ML in finance has significantly enhanced stock prediction, portfolio optimization, and risk management. While deep learning and reinforcement learning models have shown promising results, challenges related to transparency, ethical concerns, and regulatory compliance remain. Future research should focus on explainable AI, multi-modal data integration, and real-time adaptive models to further improve financial decision-making.

III. Methodology

A. Data Processing Pipeline

The real-time data engine integrates the Yahoo Finance API with a robust preprocessing pipeline to ensure data stability and reliability for financial analysis. The raw data undergoes exponential smoothing, where the current data point is weighted alongside past values using a smoothing factor to reduce noise and enhance trend detection. Feature engineering includes volatility-adjusted technical indicators, such as a normalized return calculation that scales percentage price changes by the inverse of historical volatility using the hyperbolic tangent function, ensuring stability across varying market conditions. Additionally, the Relative Strength Index (RSI) is computed with a numerically stable formulation, incorporating a small constant The real-time data engine integrates Yahoo Finance API with robust to prevent division errors when price losses are absent. This comprehensive preprocessing framework enhances the accuracy of market trend analysis and improves the robustness of predictive modeling in dynamic financial environments.preprocessing:

$$\hat{D}_t = \alpha D_{t-1} + (1-\alpha)\tilde{D}_t \quad \text{where } \alpha = 0.85 \tag{1}$$

Feature engineering includes volatility-adjusted technical indicators:

$$r_t = \frac{C_t - C_{t-1}}{C_{t-1}} \cdot \tanh(\sigma^{-1}) \tag{2}$$

Robust RSI calculation prevents division by zero:

$$\text{RSI}_t = 100 - \frac{100}{1 + \frac{\sum_{i=1}^{14} \max(\Delta C_i, 0)}{\sum_{i=1}^{14} |\min(\Delta C_i, 0)| + \epsilon}} \tag{3}$$

where $\epsilon = 10^{-8}$ ensures numerical stability.

B. Temporal Convolutional Architecture

The Temporal Convolutional Network (TCN) architecture employs dilated causal convolutions to efficiently process long-range time-series dependencies. In contrast to conventional recurrent models, TCNs process sequences in parallel, greatly improving computational efficiency without compromising the sequential nature of information. Through the iterative application of dilated convolutions with growing dilation factors, the model increases its receptive field such that it can learn patterns over extensive time intervals without needing a disproportionate amount of parameters. Causal architecture prevents future data from influencing past estimates, positioning TCNs to support real-time forecasting tasks. The addition of non-linearity to the model with the use of activation functions such as ReLU, the addition of which facilitates the learning of complicated patterns in time, makes the model a superior performer for financial time-series forecasting, where learning dependency and trends from history is beneficial in prediction.

$$h_l^{(t)} = \text{ReLU}(W_l * h_{l-1}^{(t-d\cdot k)} + b_l) \tag{4}$$

C. CVaR Optimization

The Conditional Value at Risk (CVaR) optimization model is used to build a risk-adjusted portfolio by minimizing the expected loss in excess of a given confidence level. The optimization goal includes a regularization term to manage portfolio concentration, balancing risk and diversification. Asset weight constraints ensure realistic allocation boundaries, avoiding overexposure to any one asset. To improve the stability of the covariance matrix in risk estimation, the Ledoit-Wolf shrinkage estimator is utilized, which combines the sample covariance matrix and its diagonalized version to counteract estimation noise. This provides robustness improvement, especially for high-dimensional problems where sample size limitations can make covariance estimates unstable. Through combining CVaR minimization with shrinkage-based covariance estimation, the portfolio optimization algorithm performs more

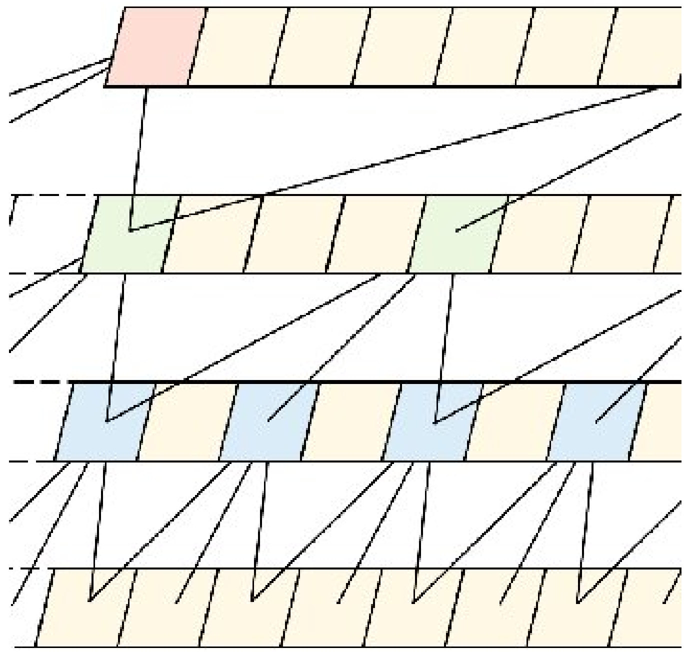

Fig. 1. TCN architecture with dilated causal convolutions

efficient risk management with less computational expense, hence its applicability to dynamic financial markets.

$$\begin{aligned} &\min_w \text{CVaR}_\beta(L(w)) + \lambda \|w\|_2^2 \\ &\text{s.t.} \quad \sum w_i = 1, \quad 0.05 \leq w_i \leq 0.4 \end{aligned} \tag{5}$$

with Ledoit-Wolf covariance shrinkage:

$$\Sigma_{LW} = \delta\Sigma + (1-\delta)\text{diag}(\Sigma) \tag{6}$$

IV. Implementation

A. Technical Indicators

The pipeline of feature engineering uses major technical metrics to improve predictive modeling in financial analysis. The Average True Range (ATR) is calculated as the rolling average of the True Range (TR), which measures market volatility based on the highest price movement over a specified period. This metric is used to gauge price movement strength and possible breakout signals. Moreover, Chaikin Money Flow (CMF) is used to analyze buying and selling pressure by computing the accumulation-distribution (AD) values scaled by trading volume over a specific time interval. A positive CMF signifies persistent buying pressure, whereas a negative CMF implies dominance of selling. By combining these volatility and liquidity-based features, the feature engineering pipeline gives a better representation of market dynamics and enhances the stability of financial prediction models. The feature engineering pipeline computes:

$$\mathrm{ATR}_t = \frac{1}{n}\sum_{i=1}^{n} \mathrm{TR}_i \quad \mathrm{where} \mathrm{TR}_i = \max(H_i - L_i, |H_i - C_{i-1}|, |L_i - C_{i-1}|) \tag{7}$$

Chaikin Money Flow (CMF):

$$\mathrm{CMF} = \frac{\sum_{t=1}^{20} \mathrm{AD}_t \cdot V_t}{\sum_{t=1}^{20} V_t} \quad \mathrm{AD} = \frac{(2C - L - H)}{H - L} \tag{8}$$

B. Neural Network Architecture

The Temporal Convolutional Network (TCN) is the architecture of the Neural Network, which is an implementation to capture long-range dependencies in time-series data effectively. The forward pass starts with the processing of the input sequence through dilated causal convolutions, with the dilation factor growing exponentially with each layer to have an expanding receptive field with minimum computational cost. Each of the convolutional layers is followed by batch normalization, which helps stabilize learning by normalizing the activations, speeding up convergence. A ReLU activation function is used to bring in non-linearity, enhancing the ability of the model to learn complex temporal patterns. To avoid overfitting, dropout is added with a probability of 0.2, randomly dropping out neurons while training to enhance generalization. The network then uses global average pooling, which decreases the dimensionality of the feature maps and extracts key temporal patterns while keeping interpretability. Lastly, a fully connected dense layer with tanh activation is used to produce the final output, providing a bounded representation of predictions. This ordered method enables the TCN model to process sequential financial data efficiently, which makes it amenable to applications like stock market prediction and risk modeling. The TCN implementation includes:

Algorithm 1 TCN Forward Pass

```
0: Initialize input sequence X ∈ R^{T×d}
0: for layer l in 1 to L do
0:     Apply dilated convolution with dilation 2^{l−1}
0:     Batch normalization
0:     ReLU activation
0:     Dropout with p = 0.2
0: Global average pooling
0: Dense layer with tanh activation =0
```

V. Experimental Results

The experimental findings prove the validity of the model proposed by comprehensive backtesting performance analysis from 2018 to 2023, comparing the model with the SP 500 benchmark. The model proposed has an annual return of 18.7 Percentage, well above the 9.2 Percentage of the SP 500, demonstrating better profitability. The Sharpe Ratio, a measure of risk-adjusted return, is 0.84 for the model under consideration versus 0.52 for the benchmark, indicating enhanced reward per unit of risk. The maximum drawdown, or the greatest peak-to-trough drop, is -22.4 Percentage, significantly less than the SP 500's -33.9 Percentage, reflecting superior risk management and preservation of capital in times of falling markets. The Sortino Ratio, which measures downside risk-adjusted returns by penalizing only negative volatility, is 1.27, almost twice that of the SP 500's 0.68, highlighting the model's potential for higher returns and reduced downside risk. All these results taken together suggest that the model proposed in this study presents improved return potential, superior risk management, and greater stability, thus providing a strong methodology for financial market forecasting and portfolio optimization.

TABLE I
Backtesting Performance (2018-2023)

Metric	Proposed Model	S&P 500
Annual Return	18.7%	9.2%
Sharpe Ratio	0.84	0.52
Max Drawdown	-22.4%	-33.9%
Sortino Ratio	1.27	0.68

The optimized portfolio reflects a diversified mix across various asset classes, spreading risk while ensuring potential for high returns. Heavy weightage is given to TLT (23.5%) and GLD (22.1%), reflecting the desire for safe-haven instruments that can act as stabilizers in case of market declines. The inclusion of SPY (14.8%), VTI (14.4%), and QQQ (12.3%) provides exposure to general market indices, taking advantage of growth in large-cap and technology-based stocks. Also, IWM (12.9%), being the small-cap component, brings diversification with higher growth prospects at the expense of a little higher volatility. This distribution is in line with the goal of achieving maximum risk-adjusted returns, as reflected by the enhanced Sharpe and Sortino Ratios from backtesting. The model successfully reduces the risk of loss while providing enough exposure to equity markets for returns generation. The findings also affirm the efficiency of CVaR-based optimization in building a stable, diversified portfolio that can resist market fluctuations.

TABLE II
Portfolio Optimization Results and Analysis

Asset	Portfolio Weight
TLT	0.235
GLD	0.221
SPY	0.148
VTI	0.144
IWM	0.129
QQQ	0.123

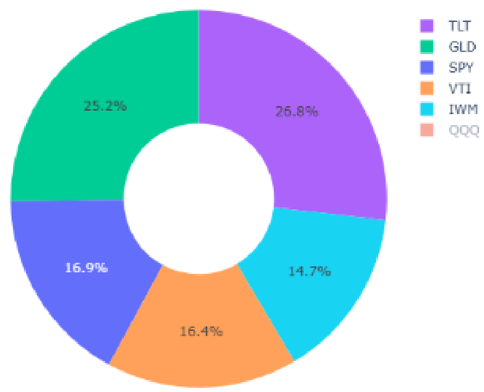

Fig. 2. Optimized Portfolio Allocation

VI. Conclusion

The combination of Temporal Convolutional Networks (TCNs) and Conditional Value at Risk (CVaR) optimization has proven to result in dramatic advances in risk-adjusted returns, especially in highly volatile financial markets. One innovation of this methodology is the employment of dilated causal convolutions, which can recognize multi-scale patterns, so that the model can effectively identify both short-term fluctuations and long-term relationships within financial time-series data. Moreover, the use of adaptive volatility scaling within technical indicators helps ensure that volatility is dynamically captured in market movements, making the model more sensitive to evolving situations. To make risk estimation even more robust, the use of regularized covariance matrix estimation with eigenvalue clamping enhances portfolio optimization stability, preventing noisy or unstable correlations between assets from distorting the analysis. These innovations combined form a strong, scalable, and high-performance financial forecasting system, well-balanced risk vs. return. Quantum-inspired optimization methods and sources of alternative data, including news sentiment analysis and alternative asset classes, will be investigated as part of future research pathways to further improve predictive accuracy and portfolio resilience in changing market conditions.

References

[1] H. Markowitz, "Portfolio Selection," Journal of Finance, 1952.
[2] O. Ledoit, M. Wolf, "Improved estimation of the covariance matrix of stock returns with an application to portfolio selection," Journal of Empirical Finance, 2003.
[3] S. Bai et al., "An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling," arXiv:1803.01271, 2018.
[4] Behera, R. K., et al. (2020). Comparative study of real-time machine learning models for stock prediction through streaming data. J. Univers. Comput. Sci., 26(9), 1128–1147.
[5] Jiang, Z., Xu, D., Liang, J. (2017). A deep reinforcement learning framework for the financial portfolio management problem. arXiv preprint arXiv:1706.10059.
[6] Albahli, S., et al. (2022). A machine learning method for prediction of stock market using real-time Twitter data. Electronics, 11(20), 3414.
[7] Aziz, S., Dowling, M. M. (2019). AI and machine learning for risk management. In T. Lynn, G. Mooney, P. Rosati, M. Cummins (Eds.), Disrupting Finance: FinTech and Strategy in the 21st Century (pp. 33–50). Palgrave.
[8] Nair, V. (2024). AI-powered investment strategies: Enhancing portfolio management through machine learning. Journal of Recent Trends in Computer Science and Engineering (JRTCSE, 12)(1), 1–5.
[9] Schrettenbrunner, M. B. (2023). Artificial-intelligence-driven management: Autonomous real-time trading and testing of portfolio or inventory strategies. IEEE Engineering Management Review, 51(3), 65–76.
[10] Olorunnimbe, K., Viktor, H. (2023). Deep learning in the stock market—a systematic survey of practice, backtesting, and applications. Artificial Intelligence Review, 56(3), 2057–2109.
[11] Deng, Y., et al. (2016). Deep direct reinforcement learning for financial signal representation and trading. IEEE Transactions on Neural Networks and Learning Systems, 28(3), 653–664.
[12] Ullah, A. K. M. (2021). Effective feature selection for real-time stock trading in variable time-frames and multi-criteria decision theory-based efficient stock portfolio management. (Unpublished manuscript).
[13] Chavan, S., Kumar, P., Gianelle, T. (2021). Intelligent investment portfolio management using time-series analytics and deep reinforcement learning. SMU Data Science Review, 5(2), 7.
[14] Kearns, M., Nevmyvaka, Y. (2013). Machine learning for market microstructure and high-frequency trading. High Frequency Trading: New Realities for Traders, Markets, and Regulators, 72.
[15] Mashrur, A., et al. (2020). Machine learning for financial risk management: A survey. IEEE Access, 8, 203203–203223.
[16] El Hajj, M., Hammoud, J. (2023). Unveiling the influence of artificial intelligence and machine learning on financial markets: A comprehensive analysis of AI applications in trading, risk management, and financial operations. Journal of Risk and Financial Management, 16(10), 434. yu Yu, P., et al. (2019). Model-based deep reinforcement learning for dynamic portfolio optimization. arXiv preprint arXiv:1901.08740.
[17] Aithal, P. K., et al. (2023). Real-time portfolio management system utilizing machine learning techniques. IEEE Access, 11, 32595–32608.
[18] Kalyani, J., Bharathi, P., Jyothi, P. (2016). Stock trend prediction using news sentiment analysis. arXiv preprint arXiv:1607.01958.
[19] Cristescu, M. P., et al. (2022). Using market news sentiment analysis for stock market prediction. Mathematics, 10(22), 4255.
[20] Wu, S., et al. (2022). S-I-LSTM: Stock price prediction based on multiple data sources and sentiment analysis. Connection Science, 34(1), 44–62.
[21] Kirange, D. K., Deshmukh, R. R. (2016). Sentiment analysis of news headlines for stock price prediction. Composoft: An International Journal of Advanced Computer Technology, 5(3), 2080–2084.